

CSE 417: Algorithms and Computational Complexity

Lecture I: Overview

Winter 2019

Larry Ruzzo



University of Washington

Computer Science & Engineering

CSE 417, Sp '14: Algorithms & Computational Complexity

[CSE Home](#)

[About Us](#) [Search](#) [Contact Info](#)

Administrative

- [FAQ](#)
- [Schedule & Reading](#)

Course Email/BBoard

- [Subscription Options](#)
- [Class List Archive](#)
- [E-mail Course Staff](#)
- [GoPost BBoard](#)

Homework

- [1: Assignment](#)
- [Electronic Turnin](#)

Lecture Notes

- [1: Intro](#)

Lecture: [EEB 037](#) ([schematic](#)) MWF 12:30- 1:20

	Office Hours	Location	Phone
Instructor: Larry Ruzzo , ruzzo@cs	tba	CSE 554	(206) 543-6298
TAs: Yanling He, hey1@cs	tba	CSE xxx	
Jianghong Shi, jhshi@cs	tba	CSE xxx	
Tianhui Shi, tianhui@cs	tba	CSE xxx	

Course Email: cse417a_sp14@uw.edu. Staff announcements and general instructor and TA are subscribed to this list. Enrolled students' Messages are automatically [archived](#).

Discussion Board: Also feel free to

Catalog Description:
Fourier T

k, lectures, etc. The [subscription options](#).

<http://courses.cs.washington.edu/417>

... algorithms for manipulating graphs and strings. Fast time and space complexity. NP-complete problems and

... Midterm, Final. Homework will be a mix of paper & pencil exercises and programming. Overall weights 55%, 15%, roughly.

Late Policy: Papers and/or electronic turnins are due at the **start** of class on the due date. 10% off for up to one day late (business day, e.g., Monday for Friday due dates); additional 20% per day thereafter.

Textbooks: [Algorithm Design](#) by [Jon Kleinberg](#) and [Eva Tardos](#). Addison Wesley, 2006. (Available from [U Book Store](#), [Amazon](#), etc.)

What you'll have to do

Homework (~55% of grade)

- Programming

 - Several small projects

- Written homework assignments

 - English exposition and pseudo-code


 - Analysis and argument as well as design

Midterm / Final Exam (~15% / 30%)

Late Policy:

Papers and/or electronic turnins are due at the *start* of class on the due date. 10% off for one day late (Monday, for Friday due dates); 20% per day thereafter.

Textbook



Algorithm Design by Jon Kleinberg and Eva Tardos. Addison Wesley, 2006.

What the course is about

Design of Algorithms

- design methods

- common or important types of problems

- analysis of algorithms - efficiency

- correctness proofs

What the course is about

Complexity, NP-completeness and intractability

solving problems in principle is not enough

algorithms must be *efficient*

some problems have *no efficient solution*

NP-complete problems

important & useful class of problems whose solutions (seemingly) cannot be found efficiently, but *can* be checked easily

Very Rough Division of Time

Algorithms (7 weeks)

Analysis of Algorithms

Basic Algorithmic Design Techniques

Graph Algorithms

Complexity & NP-completeness (3 weeks)

Check online
schedule page for
(evolving) details



University of Washington
Computer Science & Engineering

CSE 417, Wi '06: *Approximate Schedule*

[CSE Home](#)

[About Us](#) [S](#)

	Due	Lecture Topic	Reading
Week 1 1/2-1/6	M	Holiday	Ch. 1; Ch. 2
	W	Intro, Examples & Complexity	
	F	Intro, Examples & Complexity	
Week 2 1/9-1/13	M	Intro, Examples & Complexity	Ch. 3
	W	Graph Algorithms	
	F	Graph Algorithms	

Complexity Example

Cryptography (e.g., RSA, SSL in browsers)

Secret: p, q prime, say 512 bits each

Public: n which equals $p \times q$, 1024 bits

In principle

there is an algorithm that given n will find p and q :

try all $2^{512} > 1.3 \times 10^{154}$ possible p 's: kinda slow...

In practice

no fast algorithm known for this problem (on non-quantum computers)

security of RSA depends on this fact

(“quantum computing”: strongly driven by possibility of changing this)

Algorithms versus Machines

We all know about Moore's Law and the exponential improvements in hardware...


Ex: sparse linear equations over 25 years

10 orders of magnitude improvement!

Algorithms or Hardware?

25 years
progress
solving sparse
linear
systems

hardware:
4 orders of
magnitude




Algorithms or Hardware?

25 years
progress
solving sparse
linear
systems

hardware: 4
orders of
magnitude


software: 6
orders of
magnitude



Algorithms or Hardware?


The
N-Body
Problem:

in 30 years
 10^7 hardware
 10^{10} software



Algorithms or Hardware?

SAT/SMT Solvers: 1000x improvement in a dozen years



Data courtesy of Dr. Vijay Ganesh, U. Waterloo

Algorithm: definition

Procedure to accomplish a task or solve a well-specified problem

Well-specified: know what all possible inputs look like and what output looks like given them

“accomplish” via simple, well-defined steps

Ex: sorting names (via comparison)

Ex: checking for primality (via $+$, $-$, $*$, $/$, \leq)

Goals

Correctness

often subtle

Analysis

often subtle

Generality, Simplicity, 'Elegance'

Efficiency

time, memory, network bandwidth, ...


Algorithms: a sample problem

Printed circuit-board company has a robot arm that solders components to the board

Time: proportional to total distance the arm must move from initial rest position around the board and back to the initial position


For each board design, find best order to do the soldering

Printed Circuit Board



Start: Lecture 2 Jan 9

Printed Circuit Board



A Well-defined Problem

Input: Given a set S of n points in the plane

Output: The shortest cycle tour that visits each point in the set S once.

Better known as “TSP”

How might you solve it?

We could not tell you its possible runtime

Nearest Neighbor Heuristic

Start at some point p_0


Walk first to its
nearest neighbor p_1

Repeatedly walk to the nearest unvisited neighbor
 p_2 , then p_3, \dots until all points have been visited


Then walk back to p_0

heuristic: A rule of thumb, simplification, or educated guess that reduces or limits the search for solutions in domains that are difficult and poorly understood. May be good, but usually *not* guaranteed to give the best or fastest solution.


Nearest Neighbor Heuristic



An input where NN works badly




An input where NN works badly




Revised idea - Closest pairs first


Repeatedly join the closest pair of points
(s.t. result can still be part of a
single loop in the end. I.e., join
endpoints, but not points in middle,
of path segments already created.)




How does this work on our bad example?



Another bad example




Another bad example



1

$$6 + \sqrt{10} = 9.16$$

VS



8

Something that works

“Brute Force Search”:

For each of the $n! = n(n-1)(n-2)\dots 1$ orderings of the points, check the length of the cycle you get

Keep the best one

Two Notes

The two *incorrect* algorithms were “greedy”

Often very natural & tempting ideas

They make choices that look great “locally” (and never reconsider them)

When greed works, the algorithms are typically efficient

BUT: often does not work - you get boxed in

Our correct alg avoids this, but is incredibly slow

20! is so large that checking one billion orderings per second would take 2.4 billion seconds (around 70 years!)

And *growing*: $n! \sim \sqrt{2 \pi n} \cdot (n/e)^n \sim 2^{O(n \log n)}$

The Morals of the Story

Algorithms are important

- Many performance gains outstrip Moore's law

Simple problems can be hard

- Factoring, TSP

Simple ideas don't always work

- Nearest neighbor, closest pair heuristics

Simple algorithms can be very slow

- Brute-force factoring, TSP


For some problems, even the *best* algorithms are slow

Course Goals:

- formalize these ideas, and

- develop more sophisticated approaches

OUR FIELD HAS BEEN
STRUGGLING WITH THIS
PROBLEM FOR YEARS.



STRUGGLE NO MORE!
I'M HERE TO SOLVE
IT WITH *ALGORITHMS!*



SIX MONTHS LATER:

WOW, THIS PROBLEM
IS REALLY HARD.

YOU DON'T SAY.

